



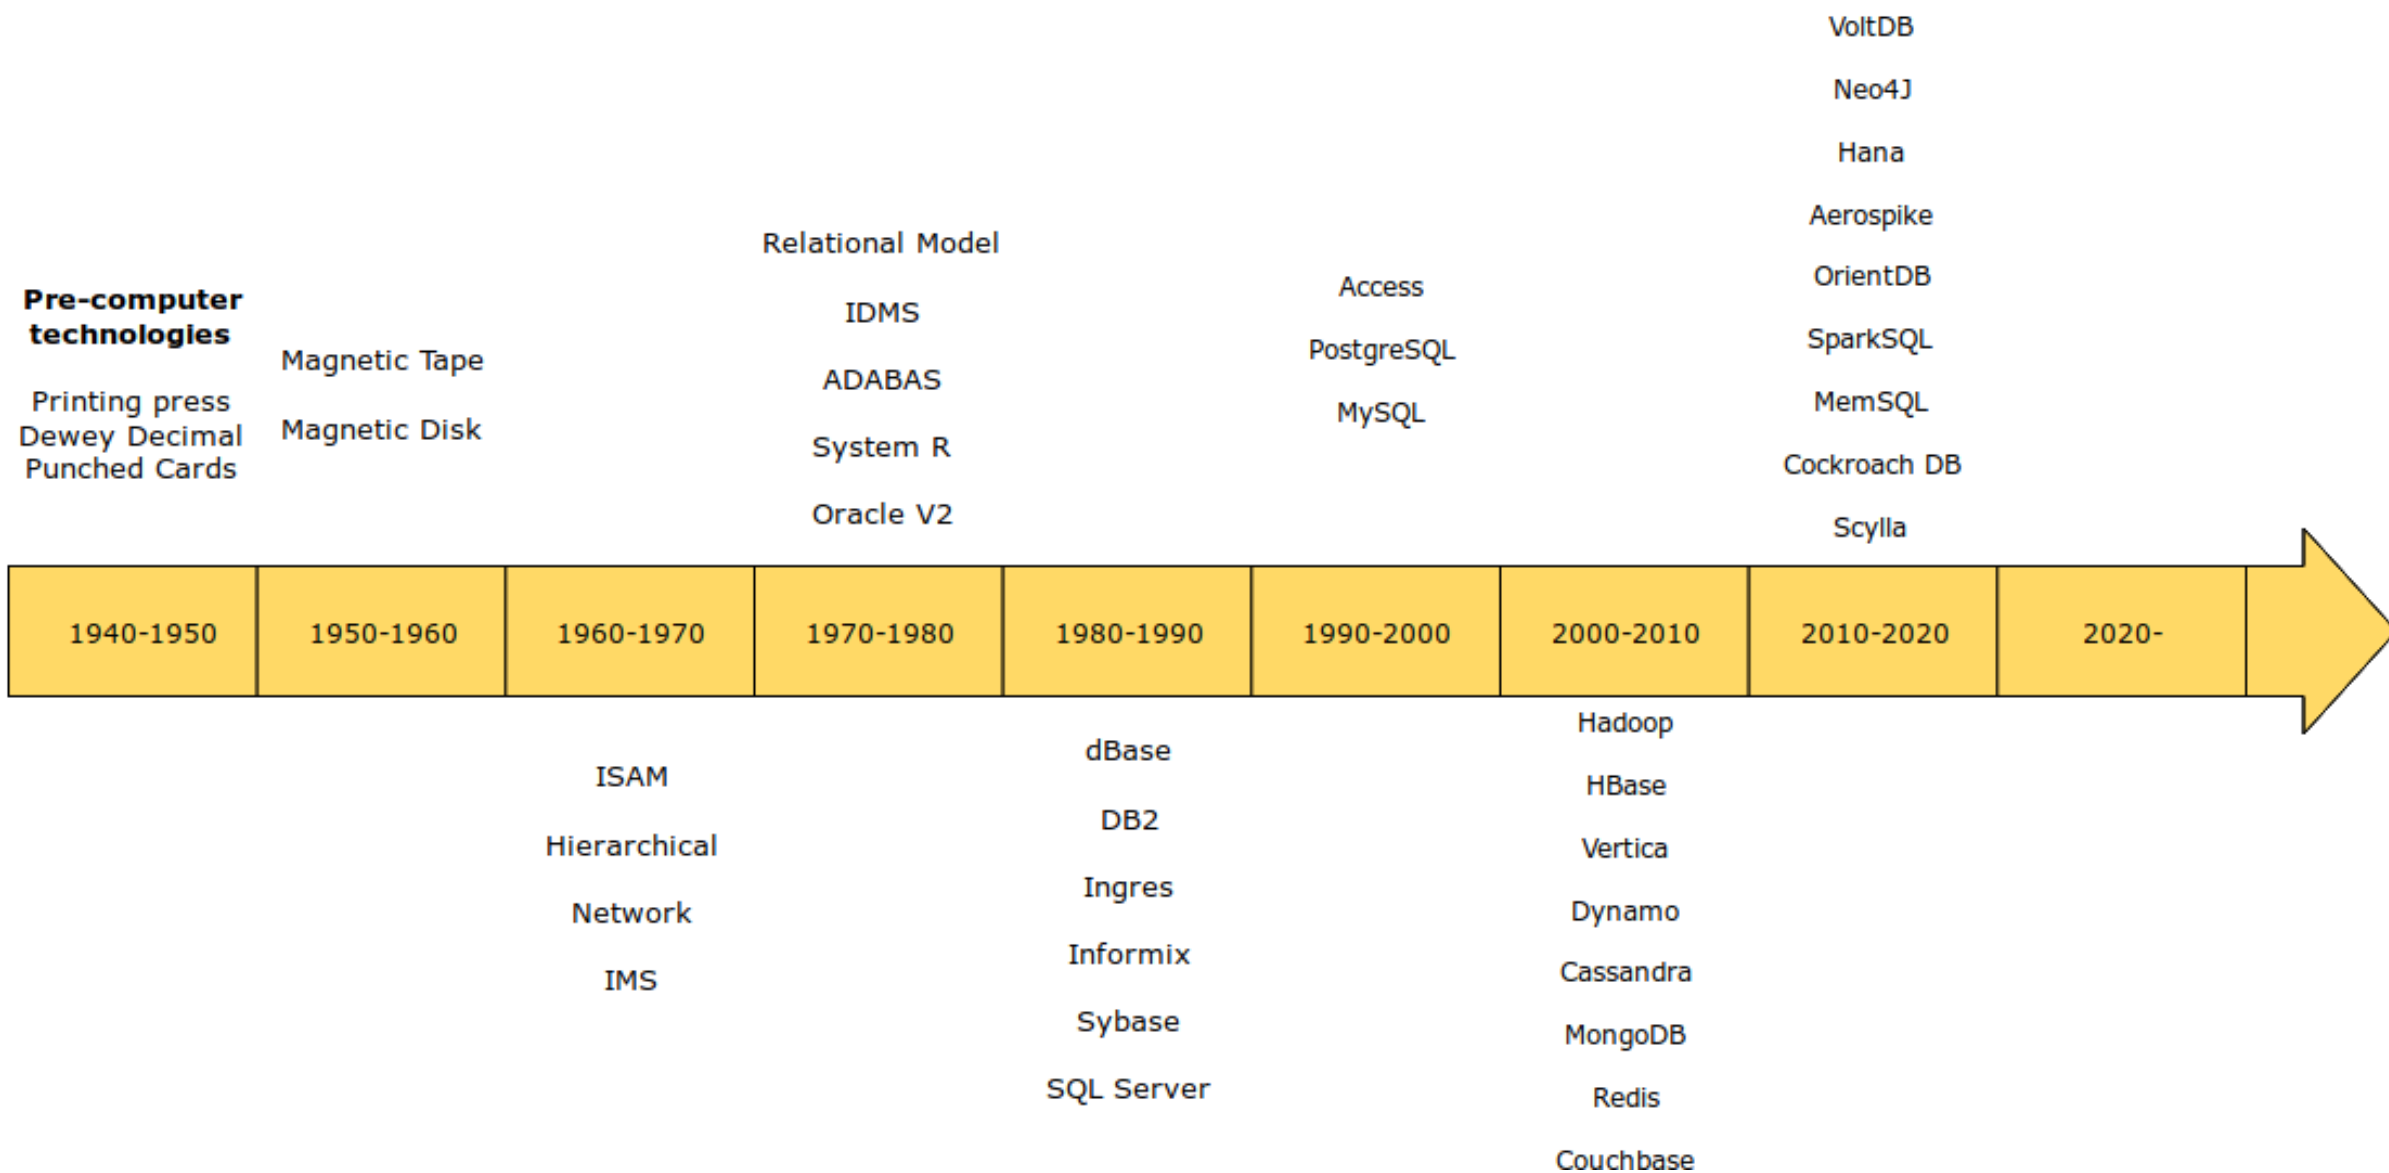
# AWS Builders Online Series

## AWS Purpose-Built Database Strategy: The Right Tool for The Right Job

Blair Layton, Business Development Manager,  
AWS

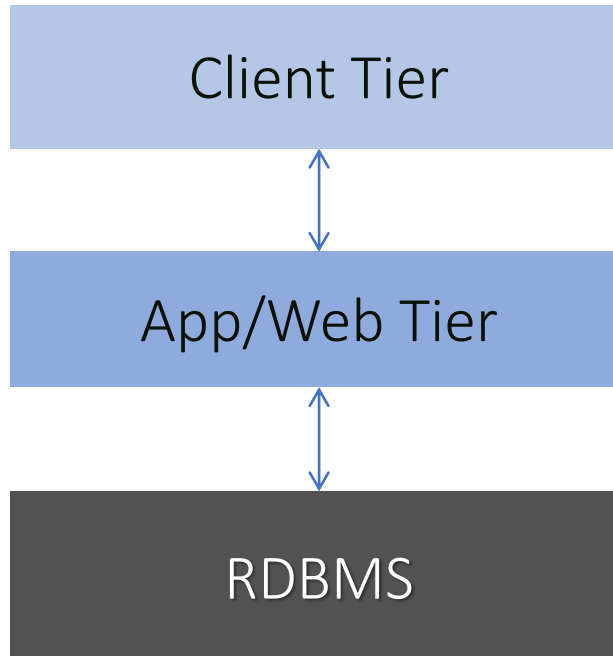


# The Database Market is Changing



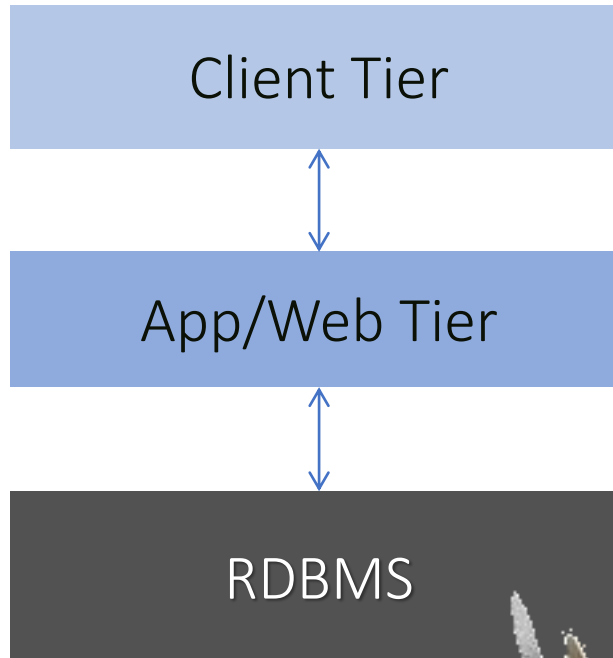
# Traditional Database Architecture

*one database for  
all workloads*



# Traditional Database Architecture

- key-value access
- complex queries
- transactions
- analytics



# Data categories and common use cases



Relational

Referential integrity, ACID transactions, schema-on-write

Lift and shift, EMR, CRM, finance



Key-value

Low-latency, key lookups with high throughput and fast ingestion of data


Real-time bidding, shopping cart, social



Document

Indexing and storing documents with support for query on any attribute

Content management, personalization, mobile



In-memory

Microseconds latency, key-based queries, and specialized data structures

Leaderboards, real-time analytics, caching



Graph

Creating and navigating data relations easily and quickly

Fraud detection, social networking, recommendation engine



Search

Indexing and searching semi-structured logs and data

Product catalog, help and FAQs, full text



Time-series

Collect, store, and process data sequenced by time

IoT applications, event tracking




Ledger

Complete, immutable, and verifiable history of all changes to application data

Systems of record, supply chain, health care, registrations, financial

# Application Architecture is Changing

# Microservices at Amazon



Service-Oriented Architecture (SOA)

Single-purpose


Connect only through APIs

Connect over HTTPS



“Microservices”





# Monolithic vs. SOA vs. Microservices



Monolithic  
Single Unit



SOA  
Coarse-grained



Microservices  
Fine-grained

# Microservices...

Eliminates any long-term commitment to a technology stack

Polyglot **ecosystem**

Polyglot **persistence**

- Decompose Databases
- **Database per microservice pattern**

Allows easy use of Canary and **Blue-Green** deployments

# Purpose-Built Databases

# Our portfolio

Broad and deep portfolio, purpose-built for builders

## Business Intelligence & Machine Learning



QuickSight



SageMaker

### Analytics



**Redshift**  
Data warehousing



**EMR**  
Hadoop + Spark



**Athena**  
Interactive analytics



**Elasticsearch Service**  
Operational Analytics



**Kinesis Data Analytics** Real time

### Databases



**Aurora**  
MySQL, PostgreSQL



**DynamoDB**  
Key value, Document



**QLDB** NEW  
Ledger Database



**RDS**  
MySQL, PostgreSQL, MariaDB,  
Oracle, SQL Server



**ElastiCache**  
Redis, Memcached



**Timestream** NEW  
Time Series



**RDS on VMware** NEW



**Neptune**  
Graph



**DocumentDB** NEW  
Document

### Blockchain



**Managed Blockchain** NEW



**Blockchain Templates**

## Data Lake



**S3/Glacier**



**Lake Formation** NEW  
Data Lakes



**Glue**  
ETL & Data Catalog

## Data Movement

Database Migration Service | Snowball | Snowmobile | Kinesis Data Firehose | Kinesis Data Streams

# AWS: Purpose-built databases



Relational



Key-value



Document



In-memory



Graph



Search



Time-series



Ledger



Amazon RDS



Amazon  
DynamoDB



Amazon  
DocumentDB



Amazon  
ElastiCache



Amazon  
Neptune



Amazon  
Elasticsearch  
Service



Amazon  
Timestream



Amazon  
Quantum  
Ledger  
Database

Aurora

Community

Commercial



# Relational Databases

# Amazon RDS

Managed relational database service with a choice of popular database engines

Amazon  
Aurora

MySQL

PostgreSQL

MariaDB


Microsoft SQL Server

ORACLE




## Easy to administer

No need to provision infrastructure, install, and maintain DB software



## Available & durable

Automatic Multi-AZ data replication; automated backup, snapshots, and failover



## Highly scalable

Scale DB compute and storage with a few clicks; minimal downtime for your application



## Fast & secure

SSD storage and guaranteed provisioned I/O; data encryption at rest and in transit

# Amazon Aurora

MySQL and PostgreSQL compatible relational database built for the cloud


Performance and availability of commercial-grade databases at 1/10th the cost

---




## Performance & scalability

5x throughput of standard MySQL and 3x of standard PostgreSQL; scale-out up to 15 read replicas




## Availability & durability

Fault-tolerant, self-healing storage; six copies of data across three AZs; continuous backup to S3



## Highly secure

Network isolation, encryption at rest/transit



## Fully managed

Managed by RDS: no hardware provisioning, software patching, setup, configuration, or backups



# Aurora: Fastest Growing Service in AWS History

NASDAQ

SAMSUNG

FINRA

bmc

URBAN AIRSHIP

EA  
ELECTRONIC ARTS™

hulu

nielsen

Pearson

ancestry

RavenPack

zynga

ZUMBA  
FITNESS

Blackboard

THOMSON REUTERS

SAFE SOFTWARE™

INRIX

Expedia®

DOW JONES

INTERCOM

Alfresco™

THE HONEST CO.  
THE HONEST CO.

FICO®

jenny  
CRAIG

AstraZeneca

CBS Interactive

workfront

gumi  
One Step Beyond.

Experian™

delight  
media & camera works

Zillow®

Lookout

NTT  
docomo

PeopleAdmin

rackspace®

FUNNY  
& DIE

SysAid

Bristol-Myers Squibb

每日新聞

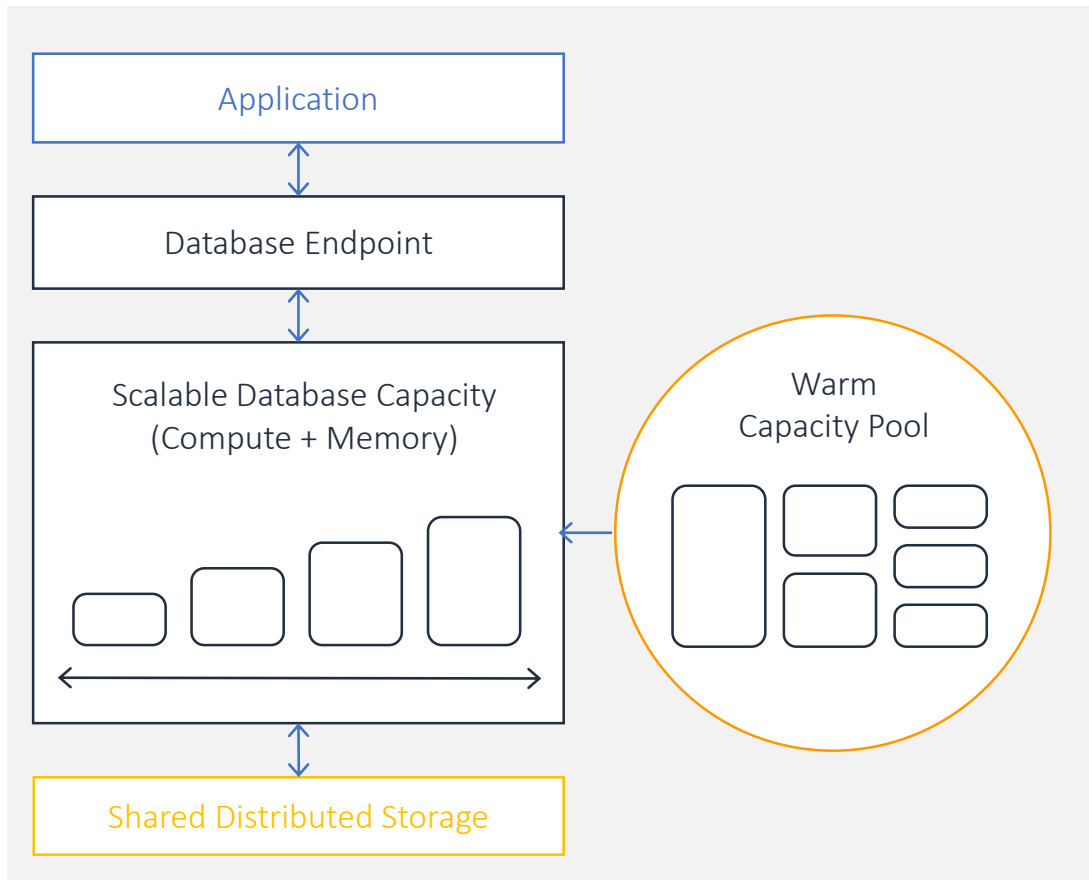
CyberAgent®

SRA OSS

TalentBin  
SOURCE THE WEB.

# Aurora Serverless

On-demand, auto-scaling database for applications with variable workloads



Starts up on demand, shuts down when not in use

Automatically scales with no instances to manage

Pay per second for the database capacity you use

# Demo

# Non-Relational Databases

# Amazon DynamoDB

We needed to adapt to power Amazon.com

## Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall and Werner Vogels

Amazon.com

### ABSTRACT

Reliability at massive scale is one of the biggest challenges we face at Amazon.com, one of the largest e-commerce operations in the world; even the slightest outage has significant financial consequences and impacts customer trust. The Amazon.com platform, which provides services for many web sites worldwide, is implemented on top of an infrastructure of tens of thousands of servers and network components located in many data centers around the world. At this scale, small and large components fail continuously and the way persistent state is managed in the face of these failures drives the reliability and scalability of the software systems.

This paper presents the design and implementation of Dynamo, a highly available key-value storage system that some of Amazon's core services use to provide an "always-on" experience. To achieve this level of availability, Dynamo sacrifices consistency under certain failure scenarios. It makes extensive use of object versioning and application-assisted conflict resolution in a manner that provides a novel interface for developers to use.

One of the lessons our organization has learned from operating Amazon's platform is that the reliability and scalability of a system is dependent on how its application state is managed. Amazon.com is a highly decentralized, loosely coupled, service oriented architecture consisting of hundreds of services. In this environment there is a particular need for storage technologies that are always available. For example, customers should be able to view and add items to their shopping cart even if disks are failing, network routes are flapping, or data centers are being destroyed by tornados. Therefore, the service responsible for managing shopping carts requires that it can always write to and read from its data store, and that its data needs to be available across multiple data centers.

Dealing with failures in an infrastructure comprised of millions of components is our standard mode of operation; there are always a small but significant number of server and network components that are failing at any given time. As such Amazon's software systems need to be constructed in a manner that treats failure handling as the normal case without impacting availability or performance.

Needed to power Amazon.com

Required massive scalability and reliability

DynamoDB designed to meet this need



# Many Applications Require Millisecond Latency at Any Scale

## Example: Amazon Prime Day 2017



**Biggest** shopping event in Amazon history

**Thousands** of Amazon teams using DynamoDB

**3.34 trillion** requests

Peaked at **12.9 million** requests per second


# Amazon DynamoDB

Fully-managed nonrelational database for any scale




## Fully managed

- Maintenance-free
- Serverless
- Auto scaling
- Backup and restore
- Global tables



## High performance


- Fast, consistent performance
- Virtually unlimited throughput
- Virtually unlimited storage



## Secure

- Encryption at rest and transit
- Fine-grained access control
- PCI, HIPAA, FIPS140-2 eligible

# DynamoDB Global Tables



Build high-performance, globally distributed applications

Low latency reads and writes to locally available tables

Multi-region redundancy and resiliency

Easy to set up and no application rewrites required



# Amazon ElastiCache

Fully-managed, Redis or Memcached compatible, low-latency, in-memory data store




## Extreme performance

In-memory data store and cache for sub-millisecond response times



## Fully managed

AWS manages all hardware and software setup, configuration, monitoring




## Easily scalable

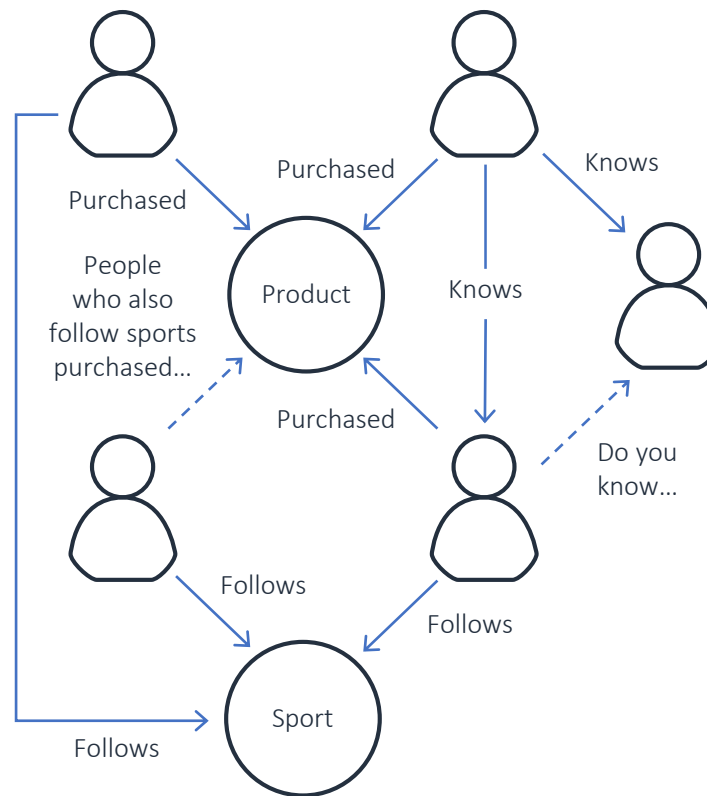
Read scaling with replicas  
Write and memory scaling with sharding  
Non-disruptive scaling

# Graph Use Cases

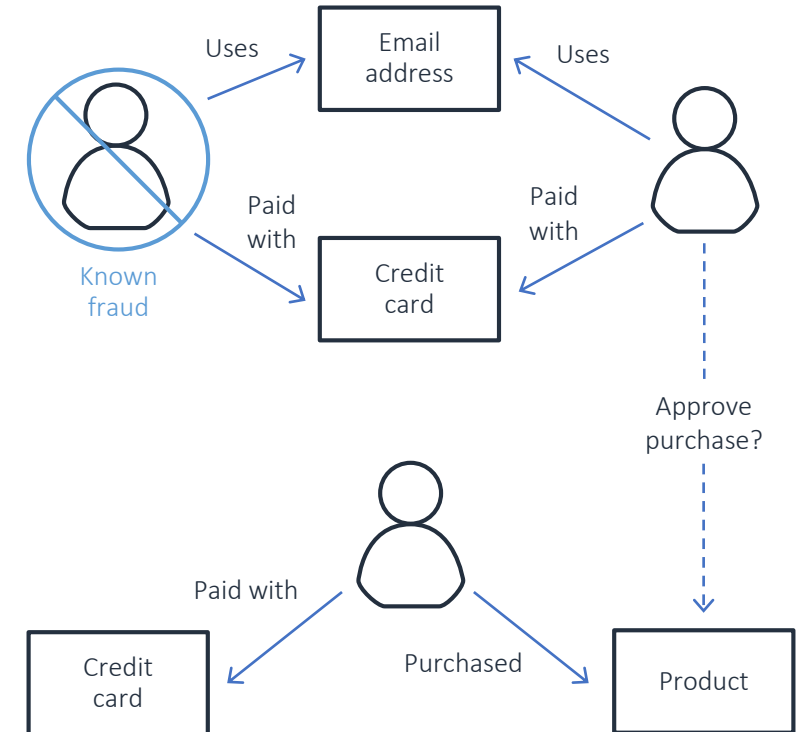
## Social news feed



## Recommendations



## Retail fraud detection



# Challenges Building Apps with Highly-Connected Data

## Relational databases



Unnatural for querying graph




Inefficient graph processing



Rigid schema inflexible for changing graphs

## Existing graph databases




Difficult to scale



Difficult to maintain high availability




Too expensive



Limited support for open standards

# Amazon Neptune

Fully managed graph database




## Open

Supports Apache TinkerPop and W3C RDF graph models



## Fast & scalable

Store billions of relationships; query with millisecond latency



## Reliable

Six replicas of your data across three AZs with full backup and restore



## Easy

Build powerful queries easily with Gremlin and SPARQL

# Amazon DocumentDB


Fast, scalable, and fully managed MongoDB-compatible database service

Fast



Millions of requests per second with millisecond latency; twice the throughput of MongoDB

Scalable



Separation of compute and storage enables both layers to scale independently; scale out to 15 read replicas in minutes

Fully managed



Managed by AWS: no hardware provisioning; auto patching, quick setup, secure, and automatic backups

MongoDB compatible



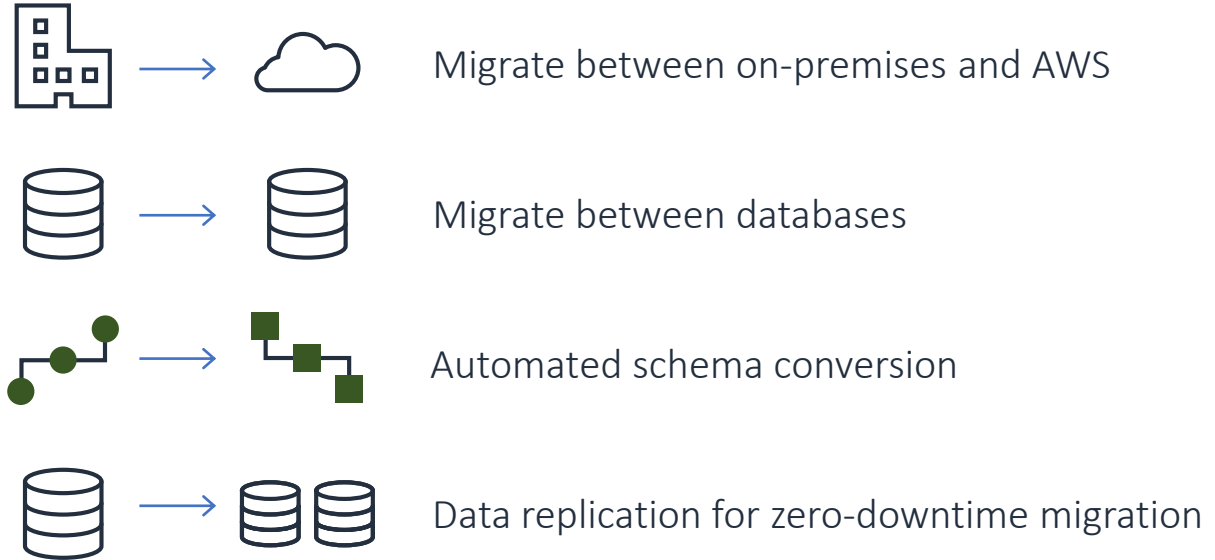
Compatible with MongoDB 3.6; use the same SDKs, tools, and applications with Amazon DocumentDB

# Database Migration Service

# AWS Database Migration Service

## Migrating Databases to AWS

100,000+  
Databases migrated



# 100,000+ Databases Migrated with DMS





# Thank you for attending AWS Builders Online Series

We hope you found it interesting! A kind reminder to **complete the survey**.  
Let us know what you thought of today's event and how we can improve the event experience for you in the future.



[aws-apac-marketing@amazon.com](mailto:aws-apac-marketing@amazon.com)



[twitter.com/AWSCloud](https://twitter.com/AWSCloud)



[facebook.com/AmazonWebServices](https://facebook.com/AmazonWebServices)



[youtube.com/user/AmazonWebServices](https://youtube.com/user/AmazonWebServices)



[slideshare.net/AmazonWebServices](https://slideshare.net/AmazonWebServices)



[twitch.tv/aws](https://twitch.tv/aws)